

Python for High Performance Data Analytics

—— (1) Computation ——

Qiyang Hu

UCLA Office of Advanced Research Computing

April 28, 2023

About this series

- The lectures will focus on
 - *High-level* overviews.
 - Introducing **libraries** that require *minimal* efforts to boost performance.
 - Short Jupyter Notebook demos
- What can/can't expected in the series?

✓ CAN	✗ CAN'T
<ul style="list-style-type: none">● From an end users' perspective	<ul style="list-style-type: none">● From a package developers' perspective
<ul style="list-style-type: none">● A <i>BIGGER</i>-picture review on the selected 3rd-party python libraries	<ul style="list-style-type: none">● Native Python tricks (e.g. container, lazy eval, mem)● Line-by-line explanations on these library interfaces
<ul style="list-style-type: none">● Demos on specific example problems	<ul style="list-style-type: none">● Discussion on the performance of various algorithms

Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos

Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos



COMPUTATION

Single Node/GPU, SIMD

- Pypy, Numba, NumExpr
- Pythran, Cython
- F2py, ctypes



DISTRIBUTED

Multiple Nodes/Machines

- MapReduce-based: PySpark, PyFlink
- MPI-based: mpi4py, Horovod
- Joblib, Dask, Ray



PYTHON



VISUALIZATION

- Viz process for big data
- Matplotlib, Bokeh, Plotly
- Holoview and Datashader
- Traited VTK, Mayavi, Paraview



DATA ARRAYS

Single Node/GPU, SIMD


- Numpy
- Pandas, Polars
- Modin, Pandarallel, Swifter
- Dask DataFrame, Vaex

Two Big **Do-Not**'s

Don't optimize prematurely.

"The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times..."


-- Donald Knuth in "TAOCP"

- 
- Easiest to understand and explain
 - Quickest to write
 - Easiest to test and maintain
 - Most portable to migrate

Don't trust benchmarks.

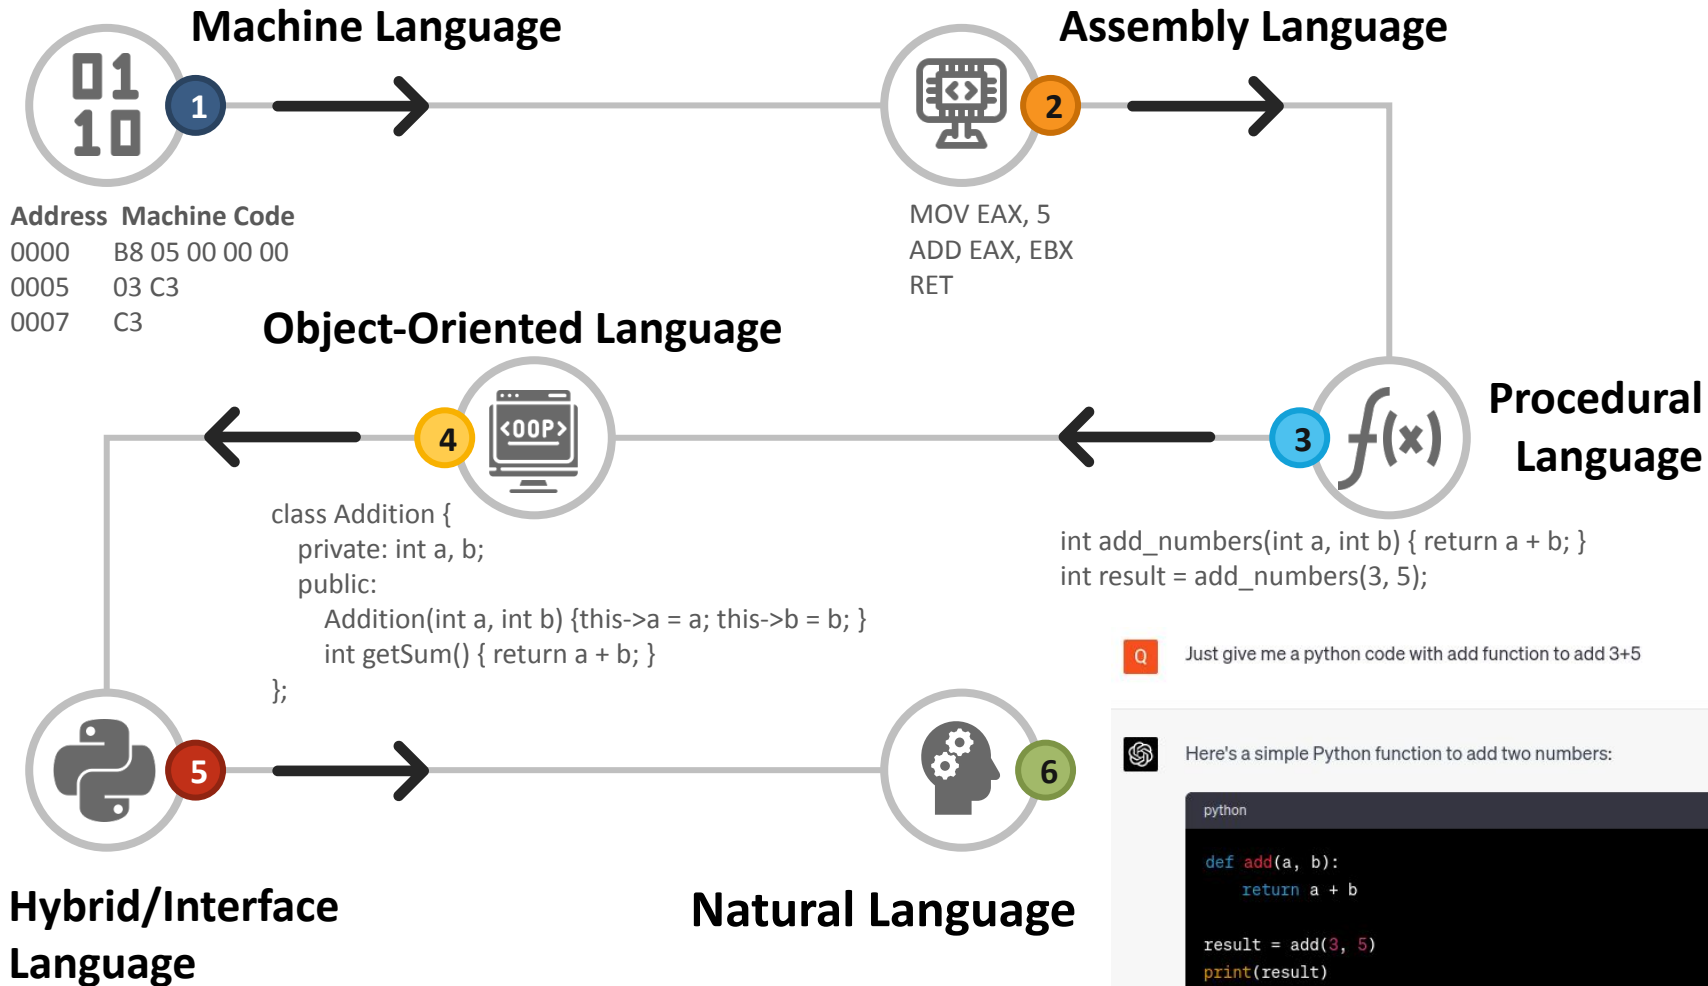
All benchmark numbers are "wrong".

- Specific hardware/OS/libraries
- In-situ running environments
- Different nature of datasets
- Sometimes very version-sensitive


- 
- Understand the mechanisms
 - Focus on the qualitative comparisons
 - Need to do your own experiments.

Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos



Q Just give me a python code with add function to add 3+5

 Here's a simple Python function to add two numbers:

```
python

def add(a, b):
    return a + b

result = add(3, 5)
print(result)
```


Why learning Python in the ChatGPT era?

Python can help us

- Understanding CS/HPC Concepts
- Providing chains of thoughts
- Nailing down the key problem quickly



Better & Efficient Prompts



Beyond the Zero/Few Shots



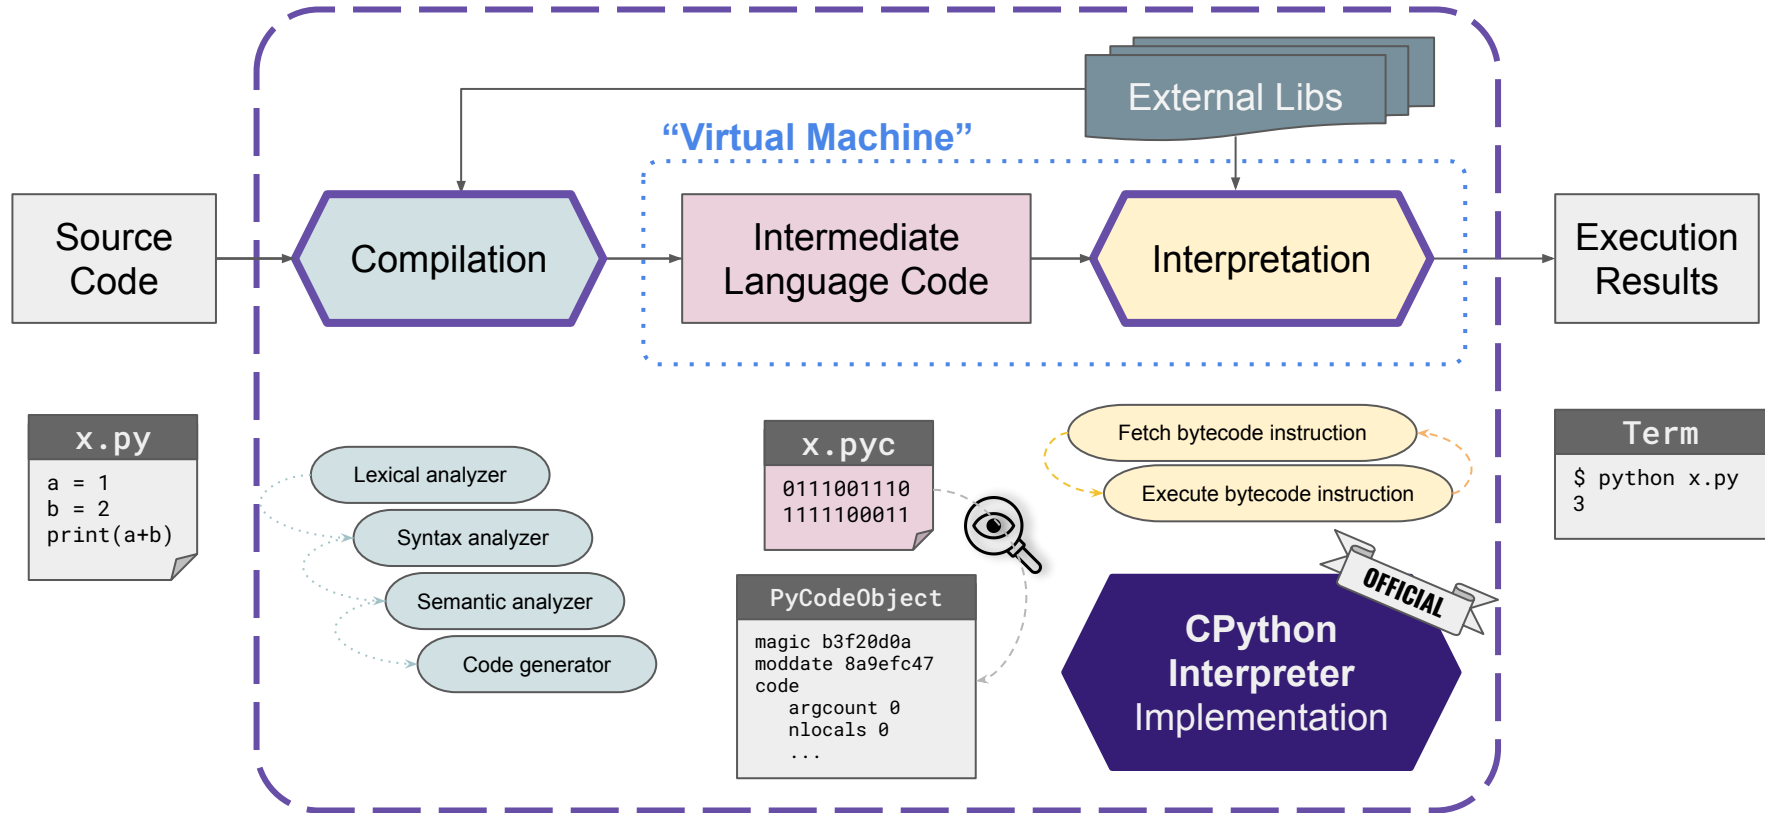
Current ChatGPT cannot do everything yet.

- Needs to fix the hallucination problem.
- Needs to distill the domain knowledge
- Needs to finetune the pre-trained model

Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- **Why Python is slow?**
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos

Seriously, what is Python?



Why Python is slow?

Python is Dynamically Typed rather than Statically Typed.

```
/* C code */  
int a = 1;  
int b = 2;  
int c = a + b;
```

```
# python code  
a = 1  
b = 2  
c = a + b
```


[Source](#)

C Addition

1. Assign `<int> 1` to `a`
2. Assign `<int> 2` to `b`
3. call `binary_add<int, int>(a, b)`
4. Assign the result to `c`

Python Addition

1. Assign `1` to `a`
 - **1a.** Set `a->PyObject_HEAD->typecode` to integer
 - **1b.** Set `a->val = 1`
2. Assign `2` to `b`
 - **2a.** Set `b->PyObject_HEAD->typecode` to integer
 - **2b.** Set `b->val = 2`
3. call `binary_add(a, b)`
 - **3a.** find typecode in `a->PyObject_HEAD`
 - **3b.** `a` is an integer; value is `a->val`
 - **3c.** find typecode in `b->PyObject_HEAD`
 - **3d.** `b` is an integer; value is `b->val`
 - **3e.** call `binary_add<int, int>(a->val, b->val)`
 - **3f.** result of this is `result`, and is an integer.
4. Create a Python object `c`
 - **4a.** set `c->PyObject_HEAD->typecode` to integer
 - **4b.** set `c->val` to `result`



The “Shannon Plan”



- A [plan](#) to make CPython faster
 - Originally proposed by Eric Snow, and Mark Shannon in 2020
 - Guido van Rossum joined and gave a talk in Python Language Summit (May 2021)
 - Based on the experience with “HotPy” and “HoyPy 2”
 - Promising 5x in 4 years, 1.5x per year
- Compatibility guarantees
 - Don’t break stable ABI compatibility
 - Don’t break limited API compatibility
 - Don’t break or slow down extreme cases
- Python 3.11 is 25% faster than 3.10
 - An adaptive, specializing bytecode interpreter and lots of optimization tweaks
 - Faster startup times and more efficient use of/communication with C
 - Will mostly benefit: CPU-intensive pure Python code

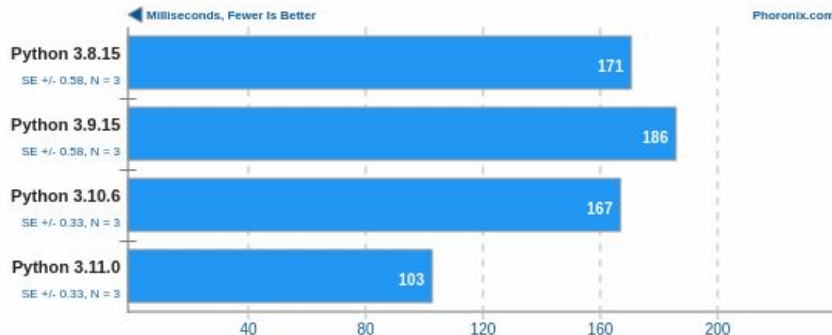
[What will 3.12 look like?](#)

PyPerformance 1.0.0

Benchmark: go



Phoronix.com




PyPerformance 1.0.0

Benchmark: float



Phoronix.com



Operation	Form	Specialization	Operation speedup (up to)	Contributor(s)
Binary operations	<code>x++x; x*x;</code> <code>x-x;</code>	Binary add, multiply and subtract for common types such as <code>int</code> , <code>float</code> , and <code>str</code> take custom fast paths for their underlying types.	10%	Mark Shannon, Dong-hee Na, Brandt Bucher, Dennis Sweeney
Subscript	<code>a[i]</code>	Subscripting container types such as <code>list</code> , <code>tuple</code> and <code>dict</code> directly index the underlying data structures. Subscripting custom <code>__getitem__</code> is also inlined similar to Inlined Python function calls .	10-25%	Irit Katriel, Mark Shannon
Store subscript	<code>a[i] = z</code>	Similar to subscripting specialization above.	10-25%	Dennis Sweeney
Calls	<code>f(arg)</code> <code>C(arg)</code>	Calls to common builtin (C) functions and types such as <code>len</code> and <code>str</code> directly call their underlying C version. This avoids going through the internal calling convention.	20%	Mark Shannon, Ken Jin
Load global variable	<code>print len</code>	The object's index in the globals/builtins namespace is cached. Loading globals and builtins require zero namespace lookups.	[1]	Mark Shannon
Load attribute	<code>o.attr</code>	Similar to loading global variables. The attribute's index inside the class/object's namespace is cached. In most cases, attribute loading will require zero namespace lookups.	[2]	Mark Shannon
Load methods for call	<code>o.meth()</code>	The actual address of the method is cached. Method loading now has no namespace lookups – even for classes with long inheritance chains.	10-20%	Ken Jin, Mark Shannon
Store attribute	<code>o.attr = z</code>	Similar to load attribute optimization.	2% in pyperformance	Mark Shannon
Unpack Sequence	<code>*seq</code>	Specialized for common containers such as <code>list</code> and <code>tuple</code> . Avoids internal calling convention.	8%	Brandt Bucher

GIL: Guilty or Gilly?

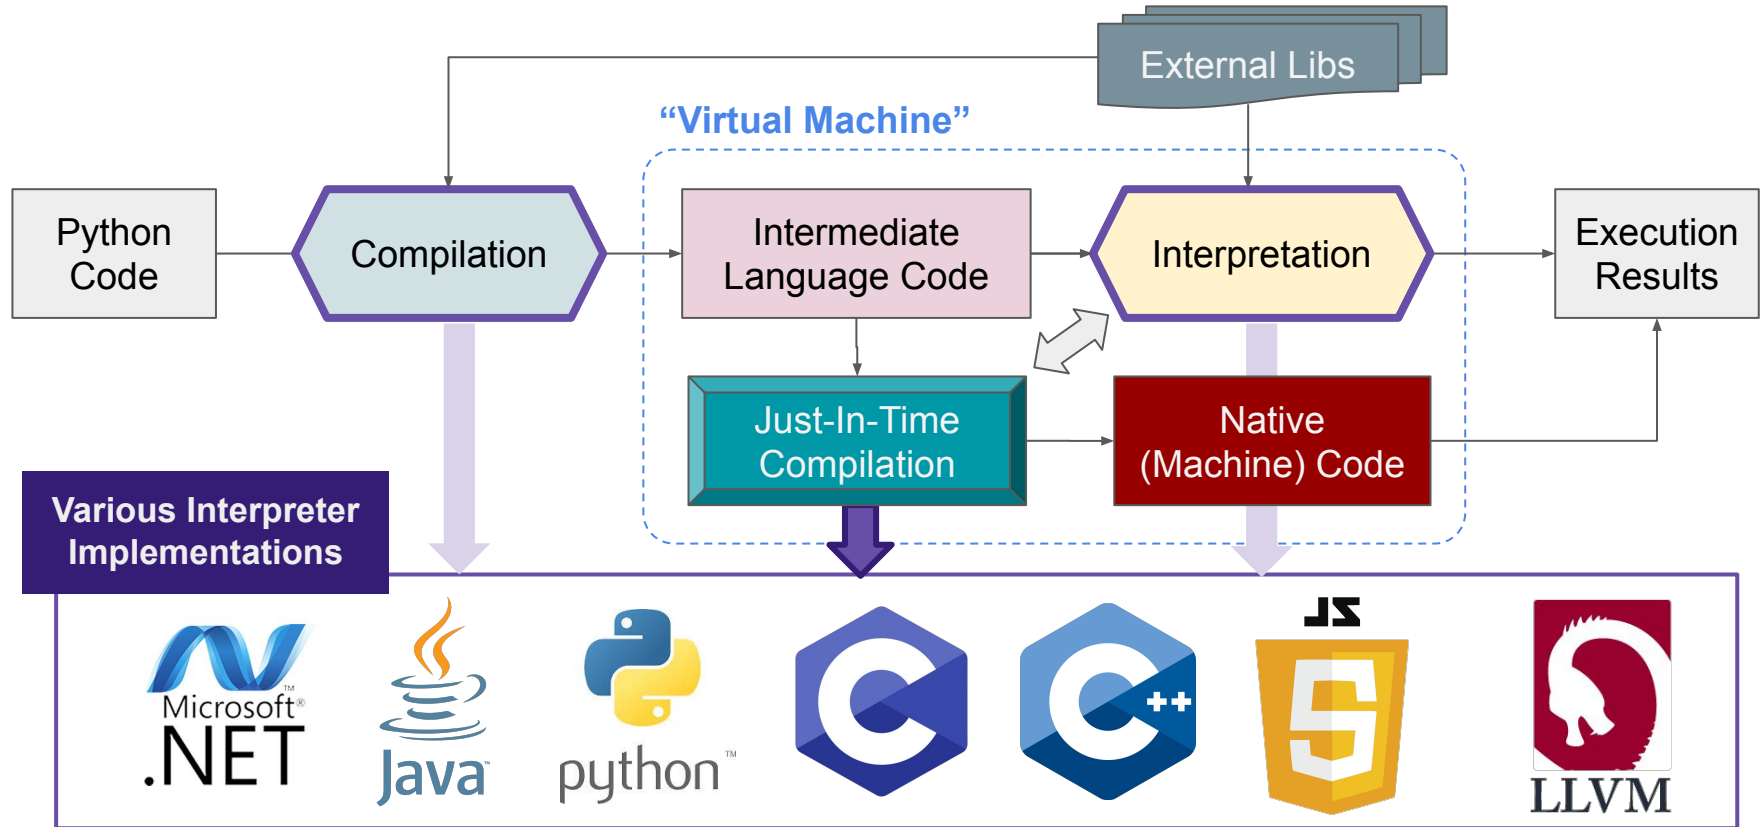


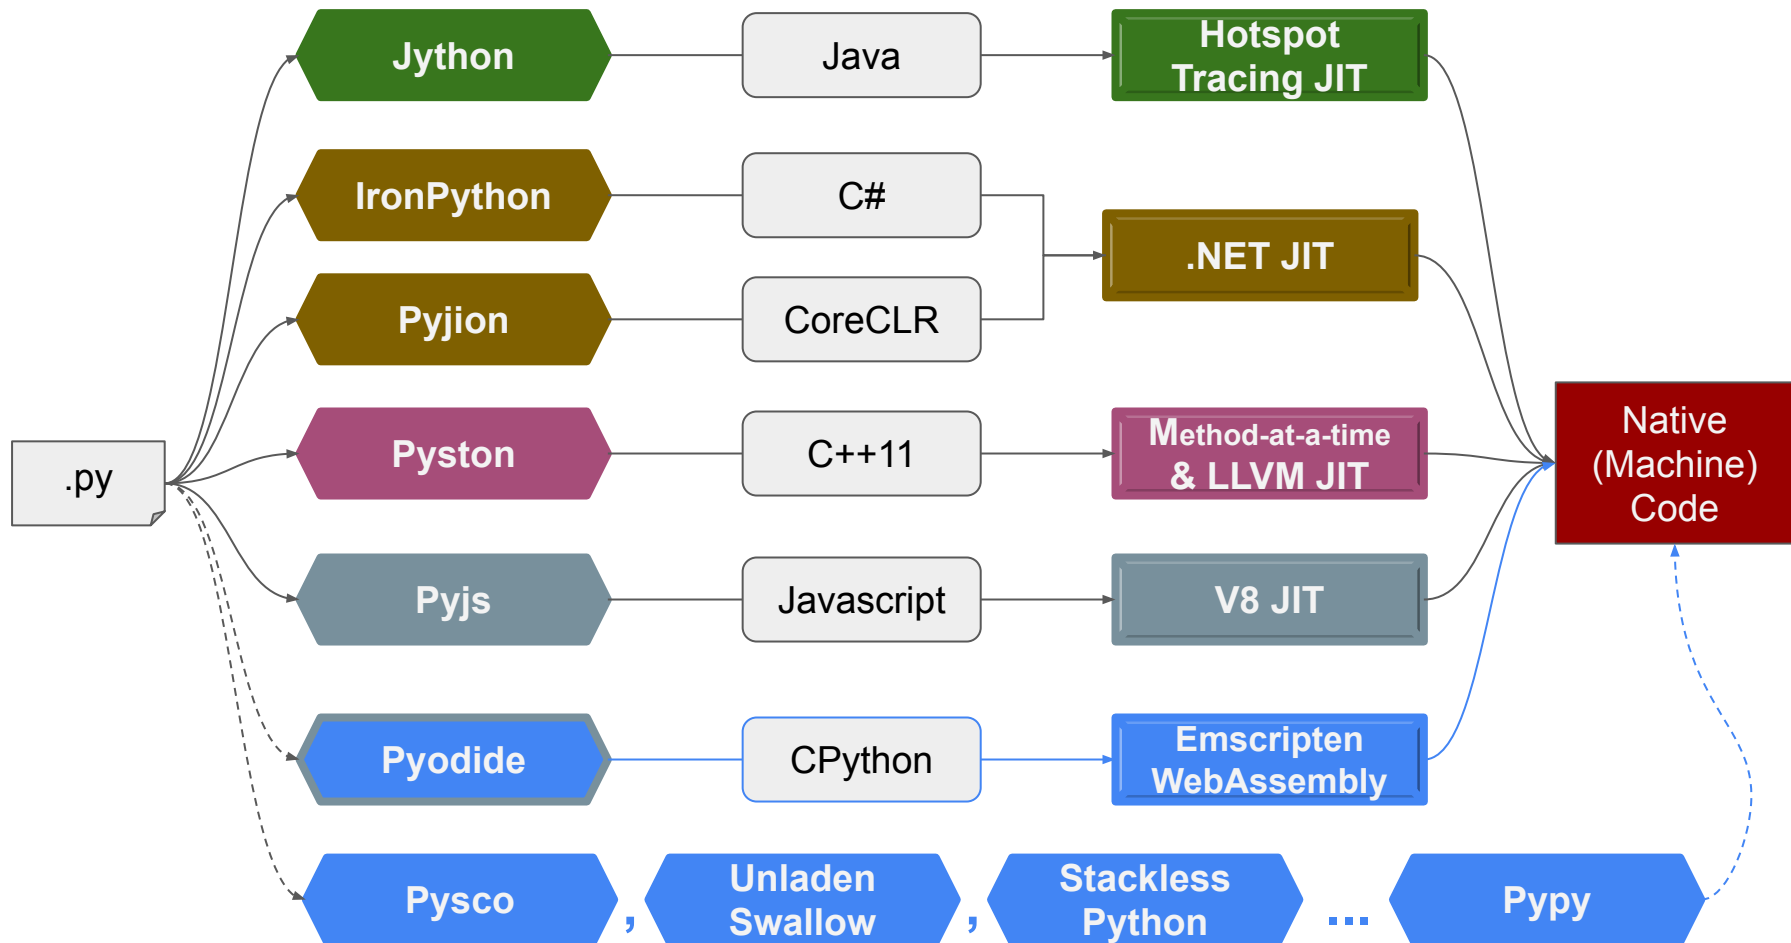
- GIL (Global Interpreter Lock)
 - A mutex (or a lock) that allows only one *thread* to hold the control of the Python interpreter.
- Why Python uses it?
 - GILs is added to the ref count variables to be kept protected from race conditions
 - GIL has performance benefits of GIL in single-threaded situation.
 - Historically Python has been around when OS did not have a concept of threads.
- Correct way to use it:
 - Multi-processing vs multi-threading:
 - Multi-threading: good for IO-intensive code, bad for CPU-intensive code
 - use multiple processes with “multiprocessing” module instead of threads
 - Consider to use Intel Distribution of Python
 - Attempts from Python community to remove the GIL from CPython:
 - [Gilectomy](#) (abandoned)
 - A new compiler flag: [nogil](#) (expected in Python 3.12)
 - Alternative Python interpreters, as GIL only with CPython
 - multiple interpreter implementations

Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos

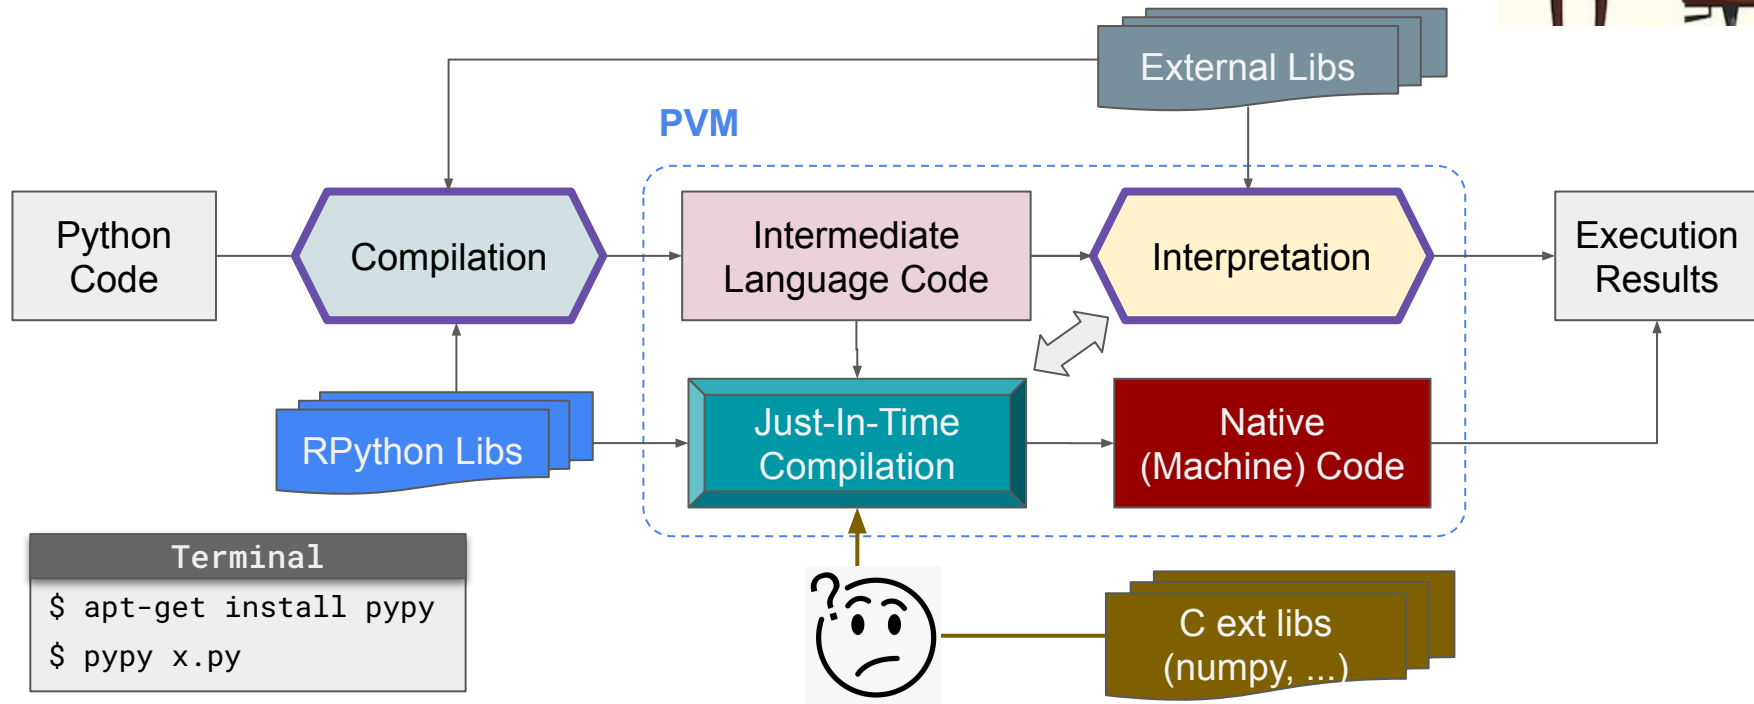
Boosting the speed by **JIT**






Pypy: using Python to interpret Python

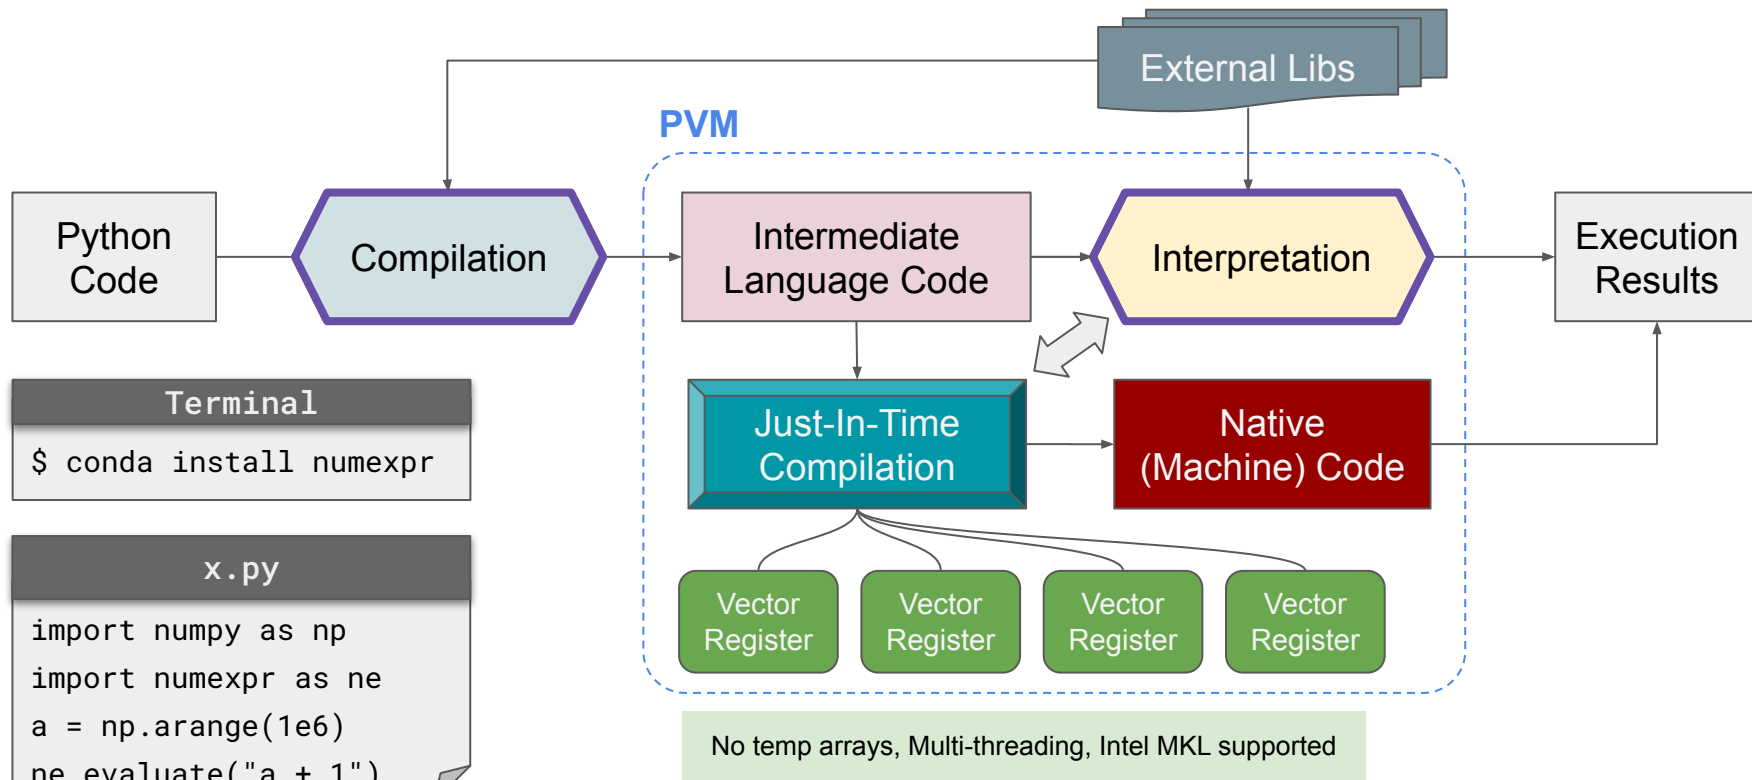
- RPython = Restricted/Reduced Python



Numba: a high-performance python JIT compiler



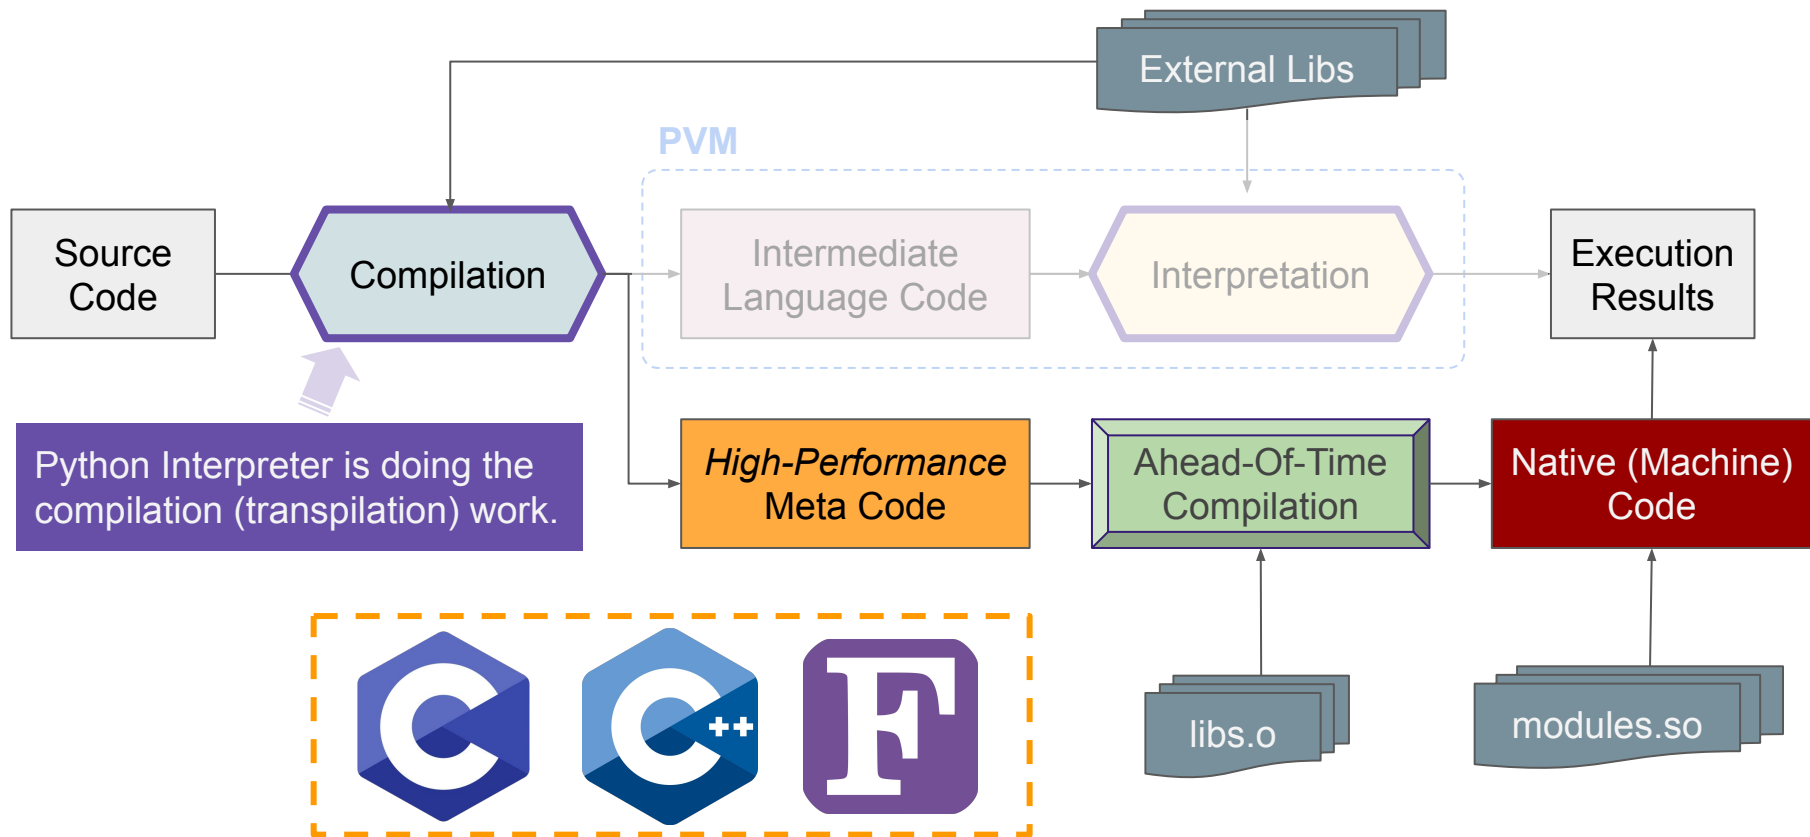
NumExpr: C-based JIT booster for numpy large arrays




Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos


Boosting the speed by **AOT** Compiler



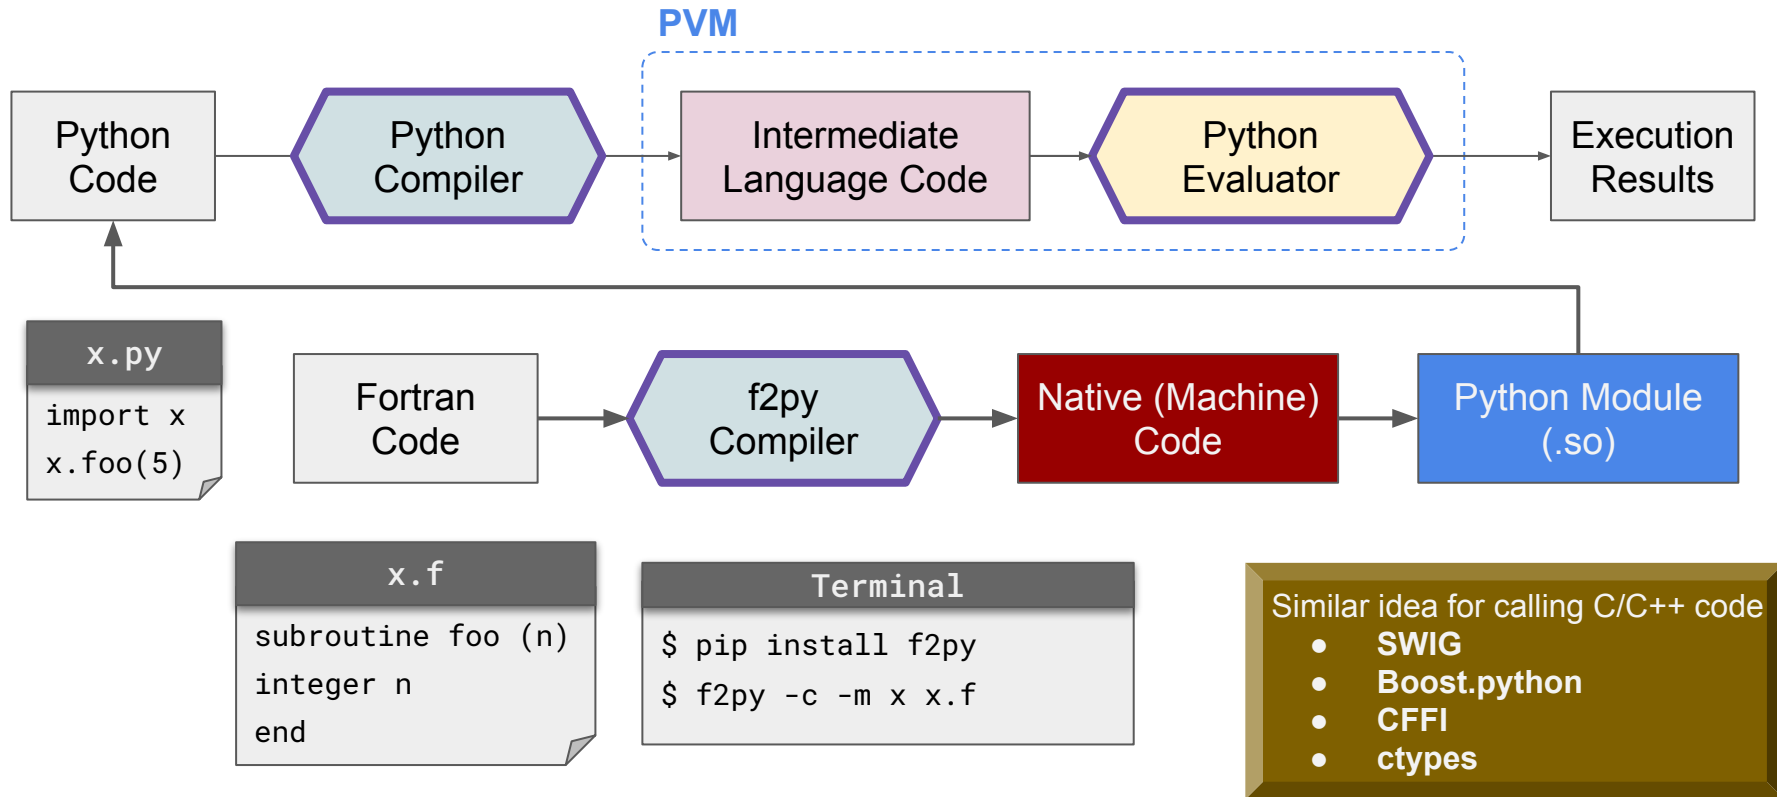
Pythran: an AOT compiler for a subset of the Python



Cython: Compiler to write C extensions for Python



f2py: wrap fortran code for use in Python



Outline for today

- What's the performance concerns for big data analytics?
- Why learning Python in the ChatGPT era?
- Why Python is slow?
- How to speed Python up using JIT?
- How to speed Python up using AOT?
- Colab demos

Hands-on Demo

bit.ly/hpdspy_01

- Code example will be running in Google Colab.
 - IPython (interpreter implementation) as Python kernel in Jupyter Notebook
 - Based on CPython, enhancing interactive features.
 - Shell prompted as `In [#]:`
 - Interacting with external files/modules by `%magic` commands
 - Some comparisons were not made in the same baseline.
 - A “maybe” game changer: **PyScript**
 - Colab comes with some installed libraries, but not all.
 - Performance benchmark was done based on array operations
 - Started with 1000 points in 3 dimensions
 - Calculate the pairwise 1000x1000 distances
 - Arrays (containers, dataframes) will be our *main* subject to discuss in the next lecture.

Key Takeaways

- Before optimization:
 - Be sure everything is working properly in a simplest way.
 - Know what's going on (data (dense/sparse?), code, algorithm, etc).
 - Profile your code and find the performance bottleneck.
- Optimization selection:
 - Consider the easiest way first (minimum changes), then the fastest.
 - Dive into the documentations
 - My personal preference order for performance optimization:
 - Scipy/Sklearn/numpy > Numba > Pythran > numexpr > Cython